

```

# TF-IDF + Klassifikation
pipeline = Pipeline([
    ('tfidf', TfidfVectorizer(max_features=5000, ngram_range=(1, 2))),
    ('clf', LogisticRegression(max_iter=1000, class_weight='balanced'))
])

# training
pipeline.fit(X_train, y_train)

# prediction
y_pred = pipeline.predict(X_test)

print("Classification Report:\n")
print(classification_report(y_test, y_pred))

print("Confusion Matrix:\n")
print(confusion_matrix(y_test, y_pred))

```

Classification Report:

	precision	recall	f1-score	support
Hate Speech	0.32	0.62	0.42	286
Neither	0.76	0.94	0.84	833
Offensive Language	0.97	0.85	0.91	3838
accuracy			0.85	4957
macro avg	0.68	0.80	0.72	4957
weighted avg	0.90	0.85	0.87	4957

Confusion Matrix:

```

[[ 177  31  78]
 [ 23 783  27]
 [ 354 212 3272]]

```

```

models = [
    ("Logistic Regression", LogisticRegression(max_iter=1000, class_weight='balanced')),
    ("Naive Bayes", MultinomialNB()),
    ("Linear SVM", LinearSVC(class_weight='balanced'))
]

results = []

for name, model in models:
    print(f"\n Testing model: {name}")

    pipe = Pipeline([
        ('tfidf', TfidfVectorizer(max_features=5000, ngram_range=(1,2))),
        ('clf', model)
    ])

    pipe.fit(X_train, y_train)
    y_pred = pipe.predict(X_test)

    acc = accuracy_score(y_test, y_pred)
    f1_macro = f1_score(y_test, y_pred, average='macro')

    print("Classification Report:")
    print(classification_report(y_test, y_pred))

    results.append((name, acc, f1_macro))

```

Testing model: Logistic Regression

Classification Report:

	precision	recall	f1-score	support
Hate Speech	0.32	0.62	0.42	286
Neither	0.76	0.94	0.84	833
Offensive Language	0.97	0.85	0.91	3838
accuracy			0.85	4957
macro avg	0.68	0.80	0.72	4957
weighted avg	0.90	0.85	0.87	4957

Testing model: Naive Bayes

Classification Report:

	precision	recall	f1-score	support
Hate Speech	1.00	0.01	0.02	286
Neither	0.90	0.49	0.64	833
Offensive Language	0.85	0.99	0.91	3838
accuracy			0.85	4957
macro avg	0.91	0.50	0.52	4957
weighted avg	0.86	0.85	0.81	4957


Testing model: Linear SVM

Classification Report:

	precision	recall	f1-score	support
Hate Speech	0.37	0.42	0.39	286
Neither	0.80	0.88	0.84	833
Offensive Language	0.94	0.91	0.93	3838
accuracy			0.88	4957
macro avg	0.70	0.74	0.72	4957
weighted avg	0.88	0.88	0.88	4957

Vergleich der Modelle:

	Model	Accuracy	F1_macro
0	Logistic Regression	0.853742	0.723605
1	Naive Bayes	0.850514	0.522986
2	Linear SVM	0.877950	0.718061



```

from transformers import TrainingArguments, Trainer
from sklearn.metrics import accuracy_score, f1_score, classification_report
from transformers import DataCollatorWithPadding

def compute_metrics(eval_pred):
    predictions, labels = eval_pred
    preds = predictions.argmax(axis=1)
    return {
        "accuracy": accuracy_score(labels, preds),
        "f1_macro": f1_score(labels, preds, average="macro"),
    }

training_args = TrainingArguments(
    output_dir="./results",
    eval_strategy="epoch",
    save_strategy="epoch",
    num_train_epochs=3,
    per_device_train_batch_size=16,
    per_device_eval_batch_size=16,
    logging_dir="./logs",
    logging_steps=50,
    load_best_model_at_end=True,
)

tokenized_datasets = tokenized_datasets.rename_column("label_id", "labels")
data_collator = DataCollatorWithPadding(tokenizer=tokenizer)

trainer = Trainer(
    model=model,
    args=training_args,
    train_dataset=tokenized_datasets["train"],
    eval_dataset=tokenized_datasets["test"],
    data_collator=data_collator,
    compute_metrics=compute_metrics,
)

```

```

trainer.train()

```

[3720/3720 3:54:12, Epoch 3/3]

Epoch	Training Loss	Validation Loss	Accuracy	F1 Macro
1	0.258300	0.261994	0.912245	0.701270
2	0.185800	0.271312	0.914666	0.717595
3	0.155900	0.292095	0.914061	0.755820

TrainOutput(global_step=3720, training_loss=0.23079304349037907, metrics={'train_runtime': 14056.9942, 'train_samples_per_second': 4.231, 'train_steps_per_second': 0.265, 'total_flos': 1969759111546368.0, 'train_loss': 0.23079304349037907, 'epoch': 3.0})

```

metrics = trainer.evaluate()
print(metrics)

# Optional: detaillierter Report
predictions = trainer.predict(tokenized_datasets["test"])
y_true = predictions.label_ids
y_pred = predictions.predictions.argmax(axis=1)
print(classification_report(y_true, y_pred, target_names=label2id.keys()))

```

```

{'eval_loss': 0.2619937062263489, 'eval_accuracy': 0.9122453096631027, 'eval_f1_macro': 0.7012698355244117, 'eval_runtime': 290.514, 'eval_samples_per_second': 17.063, 'eval_steps_per_second': 1.067, 'epoch': 3.0}

```

	precision	recall	f1-score	support
Hate Speech	0.47	0.18	0.26	286
Offensive Language	0.93	0.97	0.95	3838
Neither	0.88	0.91	0.90	833
accuracy			0.91	4957
macro avg	0.76	0.68	0.70	4957
weighted avg	0.90	0.91	0.90	4957